



Ahmed et JS
immobilier
BONJOUR

Notre EDA:

```
data = pd.read_csv("housing-train-data-6628a4723213d886993351.csv")
```

data

	Unnamed: 0	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value	ocean_
0	2072	-119.84	36.77	6.0	1853.0	473.0	1397.0	417.0	1.4817	72000.0	
1	10600	-117.80	33.68	8.0	2032.0	349.0	862.0	340.0	6.9133	274100.0	<1
2	2494	-120.19	36.60	25.0	875.0	214.0	931.0	214.0	1.5536	58300.0	
3	4284	-118.32	34.10	31.0	622.0	229.0	597.0	227.0	1.5284	200000.0	<1
4	16541	-121.23	37.79	21.0	1922.0	373.0	1130.0	372.0	4.0815	117900.0	
...	
16507	1099	-121.90	39.59	20.0	1465.0	278.0	745.0	250.0	3.0625	93800.0	
16508	18898	-122.25	38.11	49.0	2365.0	504.0	1131.0	458.0	2.6133	103100.0	N
16509	11798	-121.22	38.92	19.0	2531.0	461.0	1206.0	429.0	4.4958	192600.0	
16510	6637	-118.14	34.16	39.0	2776.0	840.0	2546.0	773.0	2.5750	153500.0	<1
16511	2575	-124.13	40.80	31.0	2152.0	462.0	1259.0	420.0	2.2478	81100.0	NEA


16512 rows × 11 columns

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 16512 entries, 0 to 16511
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            16512 non-null  int64
1   longitude             16512 non-null  float64
2   latitude              16512 non-null  float64
3   housing_median_age    16512 non-null  float64
4   total_rooms           16512 non-null  float64
5   total_bedrooms        16336 non-null  float64
6   population            16512 non-null  float64
7   households            16512 non-null  float64
8   median_income         16512 non-null  float64
9   median_house_value    16512 non-null  float64
10  ocean_proximity        16512 non-null  object
dtypes: float64(9), int64(1), object(1)
memory usage: 1.4+ MB
```

```
data['ocean_proximity'].value_counts()
```

```
<1H OCEAN    7312
INLAND       5235
NEAR OCEAN   2140
NEAR BAY     1821
ISLAND        4
Name: ocean_proximity, dtype: int64
```













```
data.isnull().sum()
```

```
Unnamed: 0          0
longitude           0
latitude            0
housing_median_age  0
total_rooms         0
total_bedrooms     176
population          0
households          0
median_income       0
median_house_value  0
ocean_proximity     0
dtype: int64
```

```
data.duplicated().sum()
```

```
0
```


```
data.hist(bins=50, figsize=(20, 15))  
plt.show()
```



```

correlation_matrix = data.corr()
plt.figure(figsize=(12, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Matrice de corrélation')
plt.show()

```



Preprocess :

```
X = df.drop(columns=["median_house_value", 'Unnamed: 0'])
y = df['median_house_value']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

numerical_pipe = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', RobustScaler())])

categorical_pipe = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False))])
```

Le choix du modèle:

```
model = XGBRegressor()

pipeline = Pipeline([('preprocessor', preprocessor),
                      ('model', model)])

pipeline.fit(X_train, y_train)

y_pred = pipeline.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("RMSE", np.sqrt(mse))
print("R^2", r2)
```

RMSE 49689.047469139674
R^2 0.8130031723227926

```
model = LinearRegression()

pipeline = Pipeline([('preprocessor', preprocessor),
                      ('model', model)])

pipeline.fit(X_train, y_train)

y_pred = pipeline.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("RMSE", np.sqrt(mse))
print("R^2", r2)
```

RMSE 67967.56067685975
R^2 0.6501222997867013

Model final et pipeline:

```
df = pd.read_csv("housing-train-data-6628a4723213d886993351.csv")

numerical_cols = ['longitude', 'latitude', 'housing_median_age',
                  'total_rooms', 'total_bedrooms', 'population', 'households',
                  'median_income']
categorical_cols = ['ocean_proximity']

X = df.drop(columns=["median_house_value", 'Unnamed: 0'])
y = df['median_house_value']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

numerical_pipe = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='median')),
    ('scaler', RobustScaler())])

categorical_pipe = Pipeline(steps=[
    ('imputer', SimpleImputer(strategy='most_frequent')),
    ('onehot', OneHotEncoder(handle_unknown='ignore', sparse_output=False))])

preprocessor = ColumnTransformer(
    transformers=[
        ('num', numerical_pipe, numerical_cols),
        ('cat', categorical_pipe, categorical_cols)])

model = KNeighborsRegressor()

pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('model', model)])

pipeline.fit(X_train, y_train)

y_pred = pipeline.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

Gridsearch:

```
param_grid = {
    'model__n_neighbors': [5, 7, 9, 12, 15, 20],
    'model__weights': ['uniform', 'distance'],
    'model__metric': ['euclidean', 'manhattan']
}
grid_pipeline = GridSearchCV(pipeline, param_grid, cv=5, scoring='neg_mean_squared_error')

grid_pipeline.fit(X_train, y_train)

y_pred = grid_pipeline.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("RMSE", np.sqrt(mse))
print("R^2", r2)
print("Meilleurs paramètres:", grid_pipeline.best_params_)
```

RMSE 58146.19510340582

R^2 0.7439318885987141

Meilleurs paramètres: {'model__metric': 'manhattan', 'model__n_neighbors': 9, 'model__weights': 'distance'}

Joblib et utilisation du modèle:

```
loaded_model = joblib.load('modele_knn.joblib')

new_data = pd.read_csv("housing-tra.csv")


X_new = new_data.drop(columns=["median_house_value"])
y_new = new_data['median_house_value']

y_pred_new = loaded_model.predict(X_new)

mse_new = mean_squared_error(y_new, y_pred_new)
r2_new = r2_score(y_new, y_pred_new)

print("Mean Squared Error (Nouveau jeu de données):", np.sqrt(mse_new))
print("R^2 Score (Nouveau jeu de données):", r2_new)
```

```
Mean Squared Error (Nouveau jeu de données): 26686.794469891785
R^2 Score (Nouveau jeu de données): 0.9463563312930844
```



```
longitude = -122.22
latitude = 37.87
housing_median_age = 52
total_rooms = 1627
total_bedrooms = 280
population = 564
households = 260
median_income = 3.8547
ocean_proximity = 'NEAR BAY'
```

```
import pandas as pd

nouvelle_ligne_df = pd.DataFrame({
    'longitude': [longitude],
    'latitude': [latitude],
    'housing_median_age': [housing_median_age],
    'median_income': [median_income],
    'ocean_proximity': [ocean_proximity],
    'total_rooms': [total_rooms],
    'total_bedrooms': [total_bedrooms],
    'population': [population],
    'households': [households]
})

prix_predit = model.predict(nouvelle_ligne_df)

print("Prix prédit :", prix_predit)
```

```
Prix prédit : [284301.75649987]
```



Planning:

Mercredi : EDA + PREPROCESS

JEUDI: TEST MODEL + ITERATION PREPROCESS

VENDREDI: TEST MODEL PIPELINES GRIDSEARCH:

LUNDI: JOBLIB STREAMLIT SLIDES GIT